

Updated PRD: NowWise

Voice-First AI Lifestyle Operating System Using Telegram + OpenWhispr-Inspired Local ASR

Product Name: NowWise

Internal Platform: NowWise OS

Core AI Engine: NowWise Decision Engine

Voice Layer: NowWise Voice Intelligence Layer

Open-Source Voice Foundation: OpenWhispr-inspired ASR service

Primary Channel: Telegram

Secondary Channels: WhatsApp, LINE, Mobile App later

Product Category: AI Lifestyle OS / Food Decision Engine / Habit Intelligence Platform

Version: Updated PRD v2.0

Core Principle: The user should never wait, never fill unnecessary forms, and never feel judged.

1. Updated Product Decision

The old NowWise PRD was correct in product vision, but after analyzing OpenWhispr, one major update is needed:

Voice should become the primary input path.

Earlier priority:

1. Text
2. Voice
3. Photo
4. Barcode
5. Receipt

Updated priority:

1. Voice
2. Text
3. Product name
4. Receipt
5. Photo
6. Barcode

Why?

Because in real life, the user is not sitting calmly and entering data. The user is:

- In Reliance Fresh
- At billing counter
- In restaurant
- Driving
- Walking
- Talking fast
- Having poor internet
- Not interested in filling forms

Voice solves this better than photo.

2. Final Product Positioning

NowWise is a **voice-first AI lifestyle operating system** that helps users make better food, grocery, restaurant, product, and habit decisions through Telegram.

The user simply speaks, types, or sends a photo.

NowWise automatically understands:

- What the user said
 - What the user ate
 - What the user is buying
 - What habit they want to schedule
 - Whether this is breakfast/lunch/dinner/snack
 - The real date and time
 - The user's goal
 - Past food behavior
 - Repeated sweets/snacks
 - Product intelligence
 - Recommended next action
-

3. Updated One-Line Product Definition

NowWise is a Telegram-first, voice-first AI lifestyle companion that converts everyday speech, food photos, grocery products, receipts, and habit requests into personalized real-time recommendations and long-term behavior intelligence.

4. Why PRD Needs Update After OpenWhispr

OpenWhispr changes our thinking because it shows that voice is not just transcription.

Voice can become:

- Food logging
- Product lookup
- Reminder creation
- Grocery basket analysis
- Restaurant decision support
- AI command input
- Daily reflection
- Weekly report request
- Admin productivity tool
- Internal PRD/prompt creation tool

OpenWhispr also proves that local/private voice processing is possible using local engines such as Whisper and Parakeet.

But harsh point:


We should not install the full OpenWhispr desktop app inside NowWise production backend.


OpenWhispr is a desktop application. NowWise needs a server-side Telegram voice pipeline.

So the updated PRD must say:

Use OpenWhispr as architecture reference and local ASR foundation, not as the entire NowWise product.

5. Updated Product Architecture





6. Core Product Philosophy

Old Philosophy

The user should not manually select date and time.

Updated Philosophy

The user should simply speak or send. NowWise should capture everything else.

That means:

- No manual date
- No manual time
- No manual meal type
- No manual form
- No forced photo upload
- No long onboarding

- No calorie obsession
- No shame language
- No medical diagnosis
- No selling personal data

7. Updated V1 Product Scope

V1 Must-Have

No.	Feature	Updated Decision	Why
1	Telegram bot	Keep	Lowest friction
2	Voice food logging	Upgrade to core	Voice is now primary
3	OpenWhisper-style ASR service	Add	Gives local/open-source voice foundation
4	Cloud ASR fallback	Add	Needed if local ASR fails
5	Text food logging	Keep	Works in poor network
6	Product name search	Keep	Fast supermarket flow
7	Auto timestamp	Keep	User should not select date/time
8	Auto meal type	Keep	Breakfast/lunch/snack/dinner inference
9	Habit reminders	Keep	Daily retention loop
10	Isabgol reminder	Keep	Strong concrete use case
11	Taken/Later/Skip buttons	Keep	Converts reminder into data
12	Weekly behavior report	Keep	Shows user value
13	Product intelligence database	Keep	Required for packaged product decisions
14	Offline/background photo queue	Keep	Photo must not block user
15	Admin review queue	Keep	Needed for low-confidence AI cases
16	Privacy center	Keep	Voice and health data are sensitive

8. What Should Move Out of V1

Feature	Decision	Reason
Full mobile app	Later	Telegram first validates demand
Full barcode scanner	V2	Product database dependency is high
Wearables	V2/Future	Not needed to prove usage
Doctor portal	Future	Regulatory and trust complexity
Insurance analytics	Avoid V1	Trust risk
Full fridge AI	Later	Hard and not needed now
Family mode	Future	Adds complexity
Meeting transcription	Not for NowWise V1	OpenWhisper has it, but NowWise does not need it
Voice biometrics	Avoid	Privacy risk
Emotion detection from voice	Avoid	Too sensitive

9. Updated User Flows

Flow 1: Voice Food Log

User sends Telegram voice:

I ate rice, dal, curd, and one sweet.

System steps:

1. Telegram receives voice note.
2. Backend downloads audio.
3. Audio is converted from OGG/OPUS to WAV.
4. Voice router sends it to OpenWhisper local ASR.
5. If local ASR fails, cloud fallback is used.
6. Transcript is generated.
7. Intent classifier detects food log.
8. Event engine captures timestamp.
9. Meal type is inferred.
10. Food items are extracted.
11. Sweet pattern is checked.

12. User receives recommendation.

Reply:

Logged as lunch. You had rice, dal, curd, and one sweet. Since this is your fourth sweet this week, avoid dessert tonight and add protein at dinner.

Flow 2: Voice Product Lookup

User at Reliance sends voice:

Can I buy Tony Garden sunflower seeds?

System:

1. Transcribes voice.
2. Detects product lookup.
3. Searches NowWise Product Graph.
4. Finds sunflower seeds.
5. Checks salted/unsalted uncertainty.
6. Checks user's weekly snack/sodium pattern.
7. Replies.

Reply:

You can buy it occasionally. If it is salted, keep it to a small handful. If unsalted, it is a better option. Since you already had salty snacks twice this week, choose unsalted if available.

Flow 3: Voice Reminder Creation

User sends voice:

Remind me every night at 9:30 to take Isabgol.

System:

1. Transcribes voice.
2. Detects reminder creation.
3. Extracts habit: Isabgol.
4. Extracts frequency: every night.
5. Extracts time: 9:30 PM.
6. Creates reminder.
7. Confirms with user.

Reply:

Done. I'll remind you every night at 9:30 PM to take Isabgol with enough water.

✓ Taken | ⌚ Later | ✗ Skip

Flow 4: Voice Grocery Basket

User says:

I am buying paneer, eggs, bananas, and chips.

System:

1. Detects grocery basket.
2. Categorizes items.
3. Checks protein, fruit, vegetables, snacks.
4. Suggests missing item.

Reply:

Good basket. Paneer and eggs cover protein, bananas cover fruit. Chips are okay occasionally. Add one vegetable if possible.

Flow 5: Poor Internet Photo Case

User sends photo at billing counter.

System immediately replies:

Photo saved. Internet looks slow, so I'll analyze it in the background. For faster help, you can also speak or type the product name.

Then the system processes later.

This is important because photo should not be the main dependency.

10. Updated Technical Architecture

10.1 Services

Service	Responsibility
Telegram Bot Gateway	Receives text, voice, photos, receipts
Input Router	Determines input type
Audio Fetcher	Downloads Telegram voice/audio file
Audio Converter	Converts OGG/OPUS/WebM to ASR-ready format
Voice Router	Chooses ASR provider
OpenWhispr ASR Adapter	Uses local Whisper/Parakeet style service
Cloud ASR Adapter	Fallback provider
Transcript Cleaner	Cleans raw transcript
Intent Classifier	Detects food/product/reminder/grocery/menu
Event Engine	Creates timeline events
Product Graph	Stores product/brand/nutrition/ingredients
Habit Scheduler	Creates reminders
Decision Engine	Produces user-specific recommendation
Weekly Report Engine	Produces weekly story
Admin Review Queue	Reviews failed/low-confidence cases
Privacy Cleaner	Deletes raw audio after transcription

10.2 OpenWhispr-Inspired ASR Service

The ASR service should expose an OpenAI-style endpoint:

```
POST /v1/audio/transcriptions
```

Input:

```
file: audio.ogg/audio.webm/audio.wav  
model: whisper-small / whisper-medium / parakeet
```

```
language: en-IN / hi-IN / auto
prompt: custom dictionary
```

Output:

```
{
  "text": "I ate rice, dal, curd, and one sweet."
}
```

This is useful because OpenWhisper's custom ASR shim pattern already describes a service that accepts multipart audio on `/audio/transcriptions` and returns JSON with a `text` field.

11. Updated Database Schema

11.1 users

Field	Purpose
id	User ID
telegram_user_id	Telegram identity
name	User name
timezone	Local timezone
language_preference	English/Hinglish/Hindi/Telugu
created_at	Signup date

11.2 user_profiles

Field	Purpose
user_id	Link to user
primary_goal	Digestion, weight, sugar control, energy
food_preferences	Vegetarian/non-vegetarian/etc.
allergies	Optional
reminder_preferences	Quiet hours, preferred time
privacy_mode	Normal/private/local-preferred

11.3 user_events

Field	Purpose
id	Event ID
user_id	User
event_type	food_log/product_lookup/reminder/grocery/menu
source	voice/text/photo/receipt/barcode
occurred_at	Real timestamp
timezone	User timezone
inferred_meal_type	breakfast/lunch/snack/dinner
raw_input	Original text/transcript
ai_interpretation	Structured meaning
confidence_score	AI confidence
status	processed/review/failed

11.4 voice_events

Field	Purpose
id	Voice event ID
user_id	User
telegram_message_id	Telegram source
received_at	Backend time
occurred_at	Telegram message time
audio_temp_path	Temporary file
provider_used	openwhisper_local/cloud_fallback
raw_transcript	ASR output
cleaned_transcript	Cleaned text
language	Detected language
confidence_score	ASR confidence

Field	Purpose
processing_time_ms	Latency
raw_audio_deleted_at	Privacy deletion timestamp

11.5 products

Field	Purpose
id	Product ID
brand	Brand name
product_name	Product name
category	Snacks, dairy, seeds, drinks
barcode	Optional
ingredients	Ingredient text
nutrition	JSON nutrition fields
allergens	Allergen flags
processing_level	Low/medium/high
verification_status	unverified/admin_verified
created_at	Created date

11.6 product_logs

Field	Purpose
id	Product log ID
user_id	User
product_id	Matched product
query_text	User input
source	voice/text/photo/barcode
decision	can_have/avoid/occasionally
reason	AI reason

Field	Purpose
confidence_score	Match confidence
occurred_at	Timestamp

11.7 reminders

Field	Purpose
id	Reminder ID
user_id	User
habit_name	Isabgol, water, walk
frequency	daily/weekly/custom
reminder_time	Time
timezone	User timezone
active	True/false
created_from	voice/text

11.8 reminder_events

Field	Purpose
id	Reminder occurrence
reminder_id	Parent reminder
scheduled_at	Time due
sent_at	Sent time
user_action	taken/later/skip/no_response
action_at	When user responded

11.9 behavior_patterns

Field	Purpose
id	Pattern ID

Field	Purpose
user_id	User
pattern_type	sweets/sodium/late_night/low_protein
period_start	Start date
period_end	End date
evidence	Events supporting pattern
insight	Human-readable insight
severity	low/medium/high

12. AI Modules

Module	Purpose
ASR module	Turns voice into text
Transcript cleaner	Cleans filler words
Intent classifier	Determines what user wants
Entity extractor	Extracts food/product/time/habit
Meal inference engine	Infers breakfast/lunch/snack/dinner
Product matcher	Matches packaged products
Habit parser	Creates reminders
Decision engine	Gives can-have/avoid/occasionally
Pattern detector	Finds repeated sweets/snacks/habits
Weekly story generator	Summarizes behavior
Safety guardrail	Avoids diagnosis/medical claims

13. Updated AI Decision Formula

Transcript/Input
 + User Goal
 + Current Time
 + Meal Window

+ Today's Timeline
+ Weekly Patterns
+ Product/Food Data
+ Confidence Score
= Personalized Decision

Example:

User:

Can I buy Tony Garden sunflower seeds?

System context:

- Time: 7:15 PM
- User goal: reduce sweets and improve digestion
- This week: salty snacks twice
- Product: sunflower seeds, maybe salted
- Confidence: medium

Reply:

You can buy it occasionally. If this is salted, keep it small. Unsalted is better today because your salty snacks are already high this week.

14. Updated Feature Priority

Must Build First

Rank	Feature	Why
1	Telegram bot	Distribution
2	Voice note support	Main UX
3	OpenWhisper-style local ASR	Privacy and cost control
4	Cloud ASR fallback	Reliability
5	Auto timestamp	No manual logging
6	Intent classifier	Converts speech into action
7	Food logging	Daily core use
8	Product lookup	Supermarket use case

Rank	Feature	Why
9	Reminder creation	Habit retention
10	Weekly report	Value proof
11	Admin review	Quality control
12	Privacy controls	Trust

Should Build Next

Feature	Why
Product compliance portal	Fix product data
Restaurant menu analysis	Strong use case
Grocery receipt scan	Works after billing
Product alternatives	Makes recommendation actionable
User correction flow	Improves accuracy
Daily reflection	Adds context

Later

Feature	Why Later
Barcode scanner	Needs product DB maturity
Wearables	Not needed for first proof
Full mobile app	Telegram first
Family mode	Complexity
Doctor reports	Requires trust and governance
B2B insights	Must wait until user trust is strong

Must Avoid

Feature	Why Avoid
Selling personal data	Destroys trust

Feature	Why Avoid
Disease prediction claims	Regulatory and ethical risk
Voice biometrics	Privacy risk
Emotion/stress inference from voice	Sensitive and unnecessary
Always-listening mode	Trust disaster
Insurance scoring	Not V1, dangerous positioning

15. Updated MVP

MVP Promise

Speak to NowWise on Telegram. It logs your food, checks products, creates reminders, and gives weekly behavior insights.

MVP Includes

Feature	Include
Telegram text	Yes
Telegram voice	Yes
Photo upload	Limited
OpenWhisper-inspired ASR	Yes
Cloud fallback	Yes
Food logs	Yes
Product lookup	Yes
Habit reminders	Yes
Isabgol flow	Yes
Weekly report	Yes
Admin review	Yes
Privacy controls	Yes

MVP Excludes

Feature	Exclude
Full app	Yes
Barcode	Yes
Wearables	Yes
Insurance	Yes
Doctor portal	Yes
Meeting transcription	Yes
OpenWhispr notes app	Yes

16. Updated Admin Panel

Admin panel should include:

Section	Purpose
Users	View user profile and goals
Events	See user logs by date/time
Voice events	See transcript success/failure
Low-confidence queue	Fix uncertain logs
Product database	Add/edit product records
Product corrections	Approve user corrections
Reminders	Check delivery and actions
Weekly reports	Preview generated summary
Privacy requests	Export/delete user data
ASR metrics	Provider latency, accuracy, fallback rate

17. Updated Technical Stack

Layer	Recommended Stack
Front channel	Telegram Bot API
Backend	Node.js / Next.js API / Python workers
Database	Supabase Postgres
File storage	Supabase Storage
Queue	Redis / BullMQ / Supabase queue
Audio conversion	FFmpeg
Local ASR	Whisper.cpp / Parakeet / sherpa-onnx
ASR inspiration	OpenWhisper
Cloud fallback	OpenAI / Groq / Deepgram
AI reasoning	GPT / Claude / Gemini
Admin dashboard	Next.js
Monitoring	Sentry + logs
Analytics	Postgres dashboards

18. Updated Implementation Plan

Phase 1: Voice Prototype

Task	Output
Create Telegram bot	Bot receives messages
Receive voice notes	Voice note event captured
Download Telegram audio	Audio stored temporarily
Convert audio	OGG/OPUS to WAV
Build ASR endpoint	<code>/v1/audio/transcriptions</code>
Connect local Whisper/Parakeet	Transcript generated
Reply transcript to Telegram	Proof complete

Phase 2: Intelligence Prototype

Task	Output
Add intent classifier	Food/product/reminder detection
Add auto timestamp	Event time captured
Add meal inference	Breakfast/lunch/snack/dinner
Add food entity extraction	Food list created
Add reminder parser	Habit reminder created
Add product lookup	Product recommendation
Add simple decision engine	Basic advice

Phase 3: MVP Launch

Task	Output
Add weekly report	Weekly story sent
Add user profile	Goals/preferences
Add product database	Basic product graph
Add admin review	Correct failed AI
Add privacy deletion	Raw audio cleanup
Add fallback ASR	Reliability
Test with 50 users	Real usage data

Phase 4: Product Intelligence

Task	Output
Add compliance portal	Admin product correction
Add grocery receipt OCR	Basket intelligence
Add restaurant menu analysis	Order guidance
Add pattern detection	Sweets, sodium, low protein
Add product alternatives	Better choices

19. Updated Success Metrics

Voice Metrics

Metric	Target
Voice transcription success	95%+
Voice intent accuracy	90%+
Reminder creation accuracy	95%+
Food entity extraction	85%+
Product lookup from voice	80%+
Average voice processing time	Under 3 seconds
User correction rate	Under 10%

Product Metrics

Metric	Target
First voice log within day 1	50%+
First reminder created	30%+
Logs per user per week	10+
Product lookups per user per week	3+
Weekly report open rate	60%+
Day 7 retention	35%+
Day 30 retention	20%+

20. Updated Risk Table

Risk	Severity	Mitigation
OpenWhisper too desktop-heavy	High	Use only ASR pattern, not full app
Voice accuracy poor for Indian foods	High	Custom dictionary + admin correction

Risk	Severity	Mitigation
Product DB incomplete	High	Manual correction portal
Photo uploads slow	High	Voice/text-first design
Cloud ASR cost high	Medium	Local ASR first, cloud fallback
Privacy concerns	Very High	Delete raw audio, consent center
AI gives medical claims	High	Safety guardrails
Too many features	High	Keep MVP focused
Users do not return daily	High	Reminders + weekly story

21. Updated Engineering Prompt

Use this prompt in Claude/Codex:

We are building NowWise, a Telegram-first and voice-first AI lifestyle operating system. The user sends Telegram voice notes, text, product names, photos, grocery receipts, or reminder requests. The system must automatically capture timestamp/timezone, infer event type, classify intent, extract food/product/habit entities, and reply with a personalized recommendation.

Study OpenWhisper/openwhisper as a reference for local/open-source transcription, especially its local Whisper/Parakeet support, self-hosted transcription pattern, audio processing, and provider routing. Do not embed the full Electron desktop app into NowWise backend.

Build a server-side NowWise Voice Service with: - Telegram voice note ingestion - Audio download - OGG/OPUS to WAV conversion using FFmpeg - OpenAI-style `/v1/audio/transcriptions` endpoint - Local ASR adapter using Whisper.cpp or Parakeet - Cloud fallback adapter - Transcript cleaner - Intent classifier - Event engine - Food/product/reminder extraction - Raw audio deletion - Admin low-confidence review queue

Produce folder structure, DB schema, API routes, background jobs, test plan, and first MVP implementation steps.

22. Final Updated Founder Recommendation

The old NowWise PRD was product-correct.

But the updated version should be:

Voice-first, Telegram-first, OpenWhisper-inspired, privacy-first, product-intelligence-backed.

Do not build:

Another food tracker.

Do not build:

Another photo scanner.

Do not build:

Another calorie counter.

Build:

A voice-first lifestyle decision engine that remembers the user and helps the next decision.

The user should feel:

"I just speak to NowWise. It understands my food, products, reminders, and habits. It helps me make a better choice right now."

That is the updated PRD direction.